

```
In [1]: import os
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.preprocessing import image_dataset_from_directory

# Paths
train_dir = "/kaggle/input/tomato-leaf-disease/Tomato Leaf Disease/train"
test_dir = "/kaggle/input/tomato-leaf-disease/Tomato Leaf Disease/test"
```

```
2025-09-17 10:11:26.468656: E external/local_xla/xla/stream_executor/cuda/cuda_fft.c
c:477] Unable to register cuFFT factory: Attempting to register factory for plugin c
uFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDE
RR
E0000 00:00:1758103886.645434      36 cuda_dnn.cc:8310] Unable to register cuDNN fac
tory: Attempting to register factory for plugin cuDNN when one has already been regi
stered
E0000 00:00:1758103886.695561      36 cuda_blas.cc:1418] Unable to register cuBLAS f
actory: Attempting to register factory for plugin cuBLAS when one has already been r
egistered
```

```
In [2]: # Data augmentation
data_augmentation = keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.2),
    layers.RandomZoom(0.2),
    layers.RandomContrast(0.2),
])

# Load training & test sets
img_size = (224, 224)
batch_size = 32

train_data = image_dataset_from_directory(
    train_dir,
    image_size=img_size,
    batch_size=batch_size,
    shuffle=True
)

test_data = image_dataset_from_directory(
    test_dir,
    image_size=img_size,
    batch_size=batch_size
)

# Save class names
class_names = train_data.class_names
print("Class names:", class_names)
```

```
# Prefetch for speed
train_data = train_data.prefetch(buffer_size=tf.data.AUTOTUNE)
test_data = test_data.prefetch(buffer_size=tf.data.AUTOTUNE)
```

```
I0000 00:00:1758103912.926072      36 gpu_device.cc:2022] Created device /job:localh
ost/replica:0/task:0/device:GPU:0 with 13942 MB memory: -> device: 0, name: Tesla T
4, pci bus id: 0000:00:04.0, compute capability: 7.5
I0000 00:00:1758103912.926760      36 gpu_device.cc:2022] Created device /job:localh
ost/replica:0/task:0/device:GPU:1 with 13942 MB memory: -> device: 1, name: Tesla T
4, pci bus id: 0000:00:05.0, compute capability: 7.5
Found 15064 files belonging to 10 classes.
Found 3771 files belonging to 10 classes.
Class names: ['Tomato__Bacterial_spot', 'Tomato__Early_blight', 'Tomato__Late_bli
ght', 'Tomato__Leaf_Mold', 'Tomato__Septoria_leaf_spot', 'Tomato__Spider_mites Tw
o-spotted_spider_mite', 'Tomato__Target_Spot', 'Tomato__Tomato_Yellow_Leaf_Curl_Vi
rus', 'Tomato__Tomato_mosaic_virus', 'Tomato__healthy']
```

```
In [3]: # Load MobileNetV2 with pretrained weights
base_model = MobileNetV2(weights="imagenet", include_top=False, input_shape=(224, 224, 3))

# Freeze most layers, keep last 30 trainable for fine-tuning
base_model.trainable = True
for layer in base_model.layers[:-30]:
    layer.trainable = False

# Build model
inputs = keras.Input(shape=(224, 224, 3))
x = data_augmentation(inputs)
x = keras.applications.mobilenet_v2.preprocess_input(x)
x = base_model(x, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.3)(x)
outputs = layers.Dense(len(class_names), activation="softmax")(x)

model = keras.Model(inputs, outputs)

# Compile with Low LR for fine-tuning
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-5),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

model.summary()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/m
obilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
9406464/9406464 ————— 0s 0us/step
Model: "functional_1"
```

| Layer (type) | Output Shape | Param # |
|---|---------------------------------------|-----------|
| input_layer_1 (InputLayer) | (None , 224, 224, 3) | 0 |
| sequential (Sequential) | (None , 224, 224, 3) | 0 |
| true_divide (TrueDivide) | (None , 224, 224, 3) | 0 |
| subtract (Subtract) | (None , 224, 224, 3) | 0 |
| mobilenetv2_1.00_224 (Functional) | (None , 7, 7, 1280) | 2,257,984 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None , 1280) | 0 |
| dropout (Dropout) | (None , 1280) | 0 |
| dense (Dense) | (None , 10) | 12,810 |

Total params: 2,270,794 (8.66 MB)

Trainable params: 1,539,210 (5.87 MB)

Non-trainable params: 731,584 (2.79 MB)

```
In [4]: callbacks = [
    keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True),
    keras.callbacks.ModelCheckpoint("best_model.keras", save_best_only=True)
]

history = model.fit(
    train_data,
    validation_data=test_data,
    epochs=25,
    callbacks=callbacks
)
```

Epoch 1/25

I0000 00:00:1758104011.617025 98 cuda_dnn.cc:529] Loaded cuDNN version 90300

471/471 ————— 60s 90ms/step - accuracy: 0.3313 - loss: 2.0005 - val_accuracy: 0.5924 - val_loss: 1.2252
Epoch 2/25

471/471 ————— 42s 89ms/step - accuracy: 0.7361 - loss: 0.8297 - val_accuracy: 0.7346 - val_loss: 0.7959
Epoch 3/25

471/471 ————— 41s 86ms/step - accuracy: 0.8175 - loss: 0.5784 - val_accuracy: 0.8255 - val_loss: 0.5532
Epoch 4/25

471/471 ————— 41s 88ms/step - accuracy: 0.8528 - loss: 0.4554 - val_accuracy: 0.8706 - val_loss: 0.4144
Epoch 5/25

471/471 ————— 41s 88ms/step - accuracy: 0.8784 - loss: 0.3754 - val_accuracy: 0.8921 - val_loss: 0.3441
Epoch 6/25

471/471 ————— 41s 88ms/step - accuracy: 0.8930 - loss: 0.3289 - val_accuracy: 0.9075 - val_loss: 0.2924
Epoch 7/25

471/471 ————— 41s 88ms/step - accuracy: 0.9062 - loss: 0.2873 - val_accuracy: 0.9146 - val_loss: 0.2660
Epoch 8/25

471/471 ————— 41s 87ms/step - accuracy: 0.9189 - loss: 0.2562 - val_accuracy: 0.9223 - val_loss: 0.2412
Epoch 9/25

471/471 ————— 41s 87ms/step - accuracy: 0.9240 - loss: 0.2339 - val_accuracy: 0.9247 - val_loss: 0.2340
Epoch 10/25

471/471 ————— 41s 88ms/step - accuracy: 0.9300 - loss: 0.2146 - val_accuracy: 0.9326 - val_loss: 0.2119
Epoch 11/25

471/471 ————— 41s 88ms/step - accuracy: 0.9351 - loss: 0.1939 - val_accuracy: 0.9321 - val_loss: 0.2098
Epoch 12/25

471/471 ————— 41s 88ms/step - accuracy: 0.9395 - loss: 0.1825 - val_accuracy: 0.9374 - val_loss: 0.1950
Epoch 13/25

471/471 ————— 41s 87ms/step - accuracy: 0.9409 - loss: 0.1727 - val_accuracy: 0.9281 - val_loss: 0.2105
Epoch 14/25

471/471 ————— 41s 87ms/step - accuracy: 0.9427 - loss: 0.1720 - val_accuracy: 0.9361 - val_loss: 0.1980
Epoch 15/25

471/471 ————— 41s 88ms/step - accuracy: 0.9440 - loss: 0.1655 - val_accuracy: 0.9433 - val_loss: 0.1833
Epoch 16/25

471/471 ————— 41s 87ms/step - accuracy: 0.9508 - loss: 0.1477 - val_accuracy: 0.9419 - val_loss: 0.1856
Epoch 17/25

471/471 ————— 42s 88ms/step - accuracy: 0.9538 - loss: 0.1347 - val_accuracy: 0.9451 - val_loss: 0.1766
Epoch 18/25

471/471 ————— 41s 88ms/step - accuracy: 0.9547 - loss: 0.1365 - val_accuracy: 0.9422 - val_loss: 0.1727
Epoch 19/25

471/471 ————— 41s 88ms/step - accuracy: 0.9612 - loss: 0.1141 - val_accuracy: 0.9464 - val_loss: 0.1623

Epoch 20/25

471/471 ————— 41s 86ms/step - accuracy: 0.9604 - loss: 0.1166 - val_accuracy: 0.9462 - val_loss: 0.1665

Epoch 21/25

471/471 ————— 41s 86ms/step - accuracy: 0.9612 - loss: 0.1165 - val_accuracy: 0.9483 - val_loss: 0.1625

Epoch 22/25

471/471 ————— 41s 86ms/step - accuracy: 0.9665 - loss: 0.1038 - val_accuracy: 0.9462 - val_loss: 0.1646

Epoch 23/25

471/471 ————— 41s 88ms/step - accuracy: 0.9654 - loss: 0.1105 - val_accuracy: 0.9523 - val_loss: 0.1503

Epoch 24/25

471/471 ————— 41s 87ms/step - accuracy: 0.9713 - loss: 0.0899 - val_accuracy: 0.9480 - val_loss: 0.1564

Epoch 25/25

471/471 ————— 41s 88ms/step - accuracy: 0.9697 - loss: 0.0930 - val_accuracy: 0.9520 - val_loss: 0.1462

```
In [5]: test_loss, test_acc = model.evaluate(test_data)
        print("Test Accuracy:", test_acc)
```

118/118 ————— 6s 52ms/step - accuracy: 0.9530 - loss: 0.1421
Test Accuracy: 0.9520021080970764

```
In [7]: import requests
        from PIL import Image

        # Example image URL
        url = "https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcQDAR_Ewp3xi3MFHdiXewcu"
        response = requests.get(url, stream=True)
        img = Image.open(response.raw).convert("RGB")

        # Preprocess
        img = img.resize((224, 224))
        img_array = np.array(img) / 255.0
        img_array = np.expand_dims(img_array, axis=0)

        # Predict
        pred = model.predict(img_array)
        print("Prediction:", class_names[np.argmax(pred)])
```

1/1 ————— 0s 43ms/step
Prediction: Tomato__Late_blight

```
In [8]: # Save the whole model (architecture + weights + optimizer state)
        model.save("plant_disease_model.h5")
```